

Temperature-Aware DRAM Cache Management—Relaxing Thermal Constraints in 3-D Systems

Minxuan Zhou, Andreas Prodromou, Rui Wang^{ID}, *Member, IEEE*, Hailong Yang^{ID},
Depei Qian^{ID}, and Dean Tullsen, *Fellow, IEEE*

Abstract—High bandwidth 3-D-stacked dynamic random access memory (DRAM) has been proposed to address the memory wall in modern systems, especially when it is used as a large last-level cache (LLC). However, stacking DRAM directly on top of the processor significantly impedes the efficiency of cooling, potentially causing thermal issues both in the processor and DRAM. Dynamic thermal management (DTM) based on DRAM temperature can be heavily intrusive because the normal working temperature for DRAM is lower than the processor temperature limit. This paper shows that in many cases it is better to disable hot portions of the cache rather than apply DTM and slow down the processor. Three temperature-aware cache management mechanisms are proposed to decrease the performance impact of DTM on 3-D systems. Our experiments show these techniques can improve the performance of DRAM-targeted DTM by 26.1% on average which make 3-D systems more practical for the future high-performance computing.

Index Terms—Energy management, memory management, stacking, temperature control.

I. INTRODUCTION

3-D-STACKED dynamic random access memory (DRAM) technology has been proposed to address the memory wall issues in modern memory-intensive applications because 3-D-stacked DRAM provides higher bandwidth than traditional off-chip DDRx DRAM [1], [2]. High-bandwidth memory can significantly accelerate emerging applications, including graph processing [3], neural network [4], and high-performance computing applications [5]. 3-D stacking technologies enable shorter interconnect wire lengths which can significantly reduce the circuit delay and system power dissipation [6], [7]. Therefore, the several commercialized products

have adopted 3-D DRAM technologies. For example, the Intel's Knights Landing architecture has adopted 8–16 GB of stacked “near” 3-D MCDRAM, a memory architecture similar to hybrid memory cubes (HMCs) [8]. Other examples of HPC-focused architectures that already feature 3-D-stacked DRAM include Nvidia's Tesla P100 [9] and AMD's FirePro S9300 X2 [10].

Currently, 3-D-stacked DRAM is not yet capable of entirely replacing off-chip memory due to its limited capacity and high cost. For instance, MICRON's HMC [1] provides up to 8-GB capacity per stack and the newest HBM2 [2] provides 8-GB per stack [11]. Due to its capacity limitations, 3-D-stacked and off-chip memory technologies often co-exist in a system. Such a configuration is termed hybrid memory architecture (HMA). Researchers have proposed the use of the high-bandwidth stacked memory as a large last-level cache (LLC) [12]–[15], or as part of memory (PoM) [16]–[18], where memory addresses can be mapped to off-chip or on-stack DRAM.

Stacking DRAM directly on the processor provides a significant benefit in bandwidth due to vertical through-silicon-vias (TSVs). Building a 3-D system is challenging, especially due to the thermal concerns. Our thermal simulations on a 3-D systems show that both the processor and the DRAM layers can exceed their respective nominal temperature thresholds for a significant part of execution time with conventional cooling systems. In particular, 3-D structures prevent the power-heavy processor from dissipating heat to cooling systems, which raises the temperature of the whole package. Furthermore, even though DRAM itself does not consume enough power to seriously exacerbate the thermal issues, the high temperature of the processor significantly impacts the temperature and operation of the 3-D-stacked DRAM. As a result, the recent commercial solutions packing the processor and high-bandwidth 3-D-stacked DRAM are mostly 2.5-D [2], which puts the DRAM stack by the side of the processor instead of above it, sacrificing both the latency and throughput and giving up much of the advantage of 3-D stacking. Increased die temperature reduces data retention time in DRAM, requiring higher refresh rate to avoid data corruption [19], [20]. For DDR3 DRAM, the normal operating temperature is 85 °C [21], [22]. Beyond the normal operating temperature, the capacitors discharge at a faster rate and the default memory refresh rate is no longer sufficient to guarantee data correctness. DRAM memories are designed to

Manuscript received October 3, 2018; revised January 4, 2019 and April 13, 2019; accepted June 11, 2019. Date of publication July 9, 2019; date of current version September 18, 2020. The work of R. Wang was supported by the National Key Research and Development Program of China under Grant 2018YFB0203901. The work of D. Qian was supported by NSFC under Grant 61732002. The work of D. Tullsen was supported by NSF under Grant CNS-1652925. This paper was recommended by Associate Editor C. Zhuo. (Corresponding author: Rui Wang.)

M. Zhou, A. Prodromou, and D. Tullsen are with the Department of Computer Science and Engineering, University of California at San Diego, San Diego, CA 92093 USA.

R. Wang, H. Yang, and D. Qian are with the School of Computer Science and Engineering, Beihang University, Beijing 100083, China (e-mail: wangrui@buaa.edu.cn).

Digital Object Identifier 10.1109/TCAD.2019.2927528

address the issue by increasing their refresh rate during thermal emergencies [1], [2], which in turn wastes significantly more energy [23], [24].

Therefore, addressing thermal issues becomes critical to make 3-D-stacked DRAM feasible in the future. Conventional processors maintain operating temperatures via dynamic thermal management (DTM) solutions such as dynamic voltage–frequency scaling (DVFS) [25]–[27]. DTM mechanisms typically affect performance, even in 2-D systems. In 3-D systems, the desired operating temperature of DRAM is typically much lower than that of the processor (85 °C versus 105 °C) [28]. Due to disparate temperature sensitivities between DRAM and CPU logic, more aggressive DTM is required to guarantee the temperature limits on all the layers of the stack. Our results show that such an aggressive DTM introduces significant performance degradation—from 24% to 52%. We further find that the temperature violations on DRAM layers happen more often than the processor, even though both peak and average temperatures are higher on the CPU. Thus, the conventional CPU-targeted DTM is not sufficient to protect the memory, while DRAM-targeted DTM is too aggressively throttles the CPU.

Our proposed mechanisms are guided by the three basic observations. First, the highly skewed temperature distribution that challenges 3-D systems is primarily driven by processor components with varying power density (“hot spots”), combined with their physical distance from the cooling system. Second, the hottest DRAM layers that will first exceed nominal temperature are closer to the processor, while DRAM layers closer to the cooling system remain under the operational temperatures. Third, the lower limit on DRAM temperature is not the result of being more prone to permanent damage at a high temperature, but rather than the loss of retention time [29]. Thus, we can let some DRAM areas reach higher temperatures if we are not storing data in them.

Based on these observations, we identify an opportunity to alleviate performance degradation caused by overly aggressive DTM on 3-D systems: when faced with a DRAM-based thermal emergency, the system can avoid throttling the CPU. Instead, memory accesses originally addressed to a hot region can be redirected to available cool parts of memory. In other words, we propose mechanisms that the tradeoff LLC capacity for increased CPU frequency, to improve the overall system performance. Our quantitative exploration demonstrates that temporarily disabling regions of our 3-D-stacked LLC is often the better decision. In this paper, we propose three mechanisms that implement our proposal, varying in efficiency, complexity, granularity, and overhead. Compared to a “DRAM-safe,” aggressive DTM solution, our mechanisms improve performance by 21.3%–26.1% on average. Compared to the upper bound—a hypothetical system that ignores thermal violations and never throttles the CPUs—our mechanisms come within 4.5% of the upper bound.

II. MOTIVATION

A. Temperature Distribution in 3-D Systems

In a 3-D system, the processor and 3-D-stacked DRAM chips are packaged together and vertical TSVs are used as

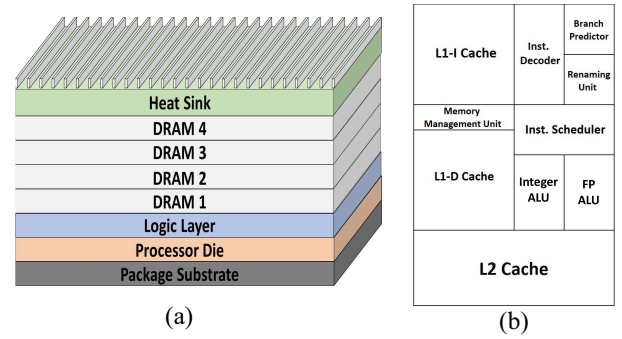


Fig. 1. (a) Memory-on-top 3-D system. (b) Single core floorplan.

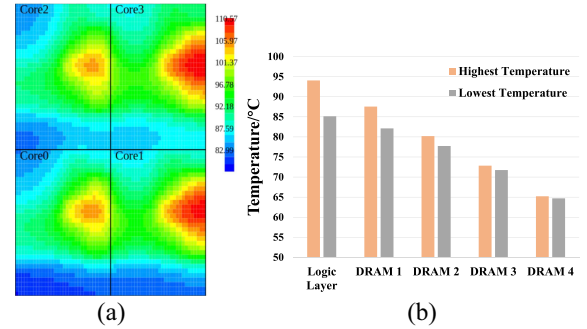


Fig. 2. (a) Horizontal temperature distribution on processor die. (b) Vertical temperature distribution on DRAM dies.

the memory bus, providing higher bandwidth than traditional off-chip memory. In order to investigate the thermal issue fairly, we assume the existence of a high-end heat sink on top of the 3-D-stacked DRAM to help the system dissipate heat. However, in a 3-D-stacked system, the distance between the underlying logic die and the heat sink becomes larger, which reduces the efficiency of the cooling system. It has been proposed to alleviate thermal issues by placing the processor die on the top of the stack, nearest the heat sink; however, this has significant manufacturing limitations as discussed by Agrawal *et al.* [30]. As a result, we utilize the much more common memory-on-top structure as shown in Fig. 1(a). The underlying processor consists of four CPU cores and the floorplan of each core is similar to that assumed by Long *et al.* [31]. Fig. 1(b) shows the layout of a single core’s components used in this paper. We introduce the details of the simulation infrastructure used in all experiments in Section V.

Fig. 2 shows the horizontal and vertical steady-state temperature distribution of our system when running *gcc* without any thermal management. The variant power density of a core’s components produces uneven horizontal temperature distribution at the base of our 3-D system. DRAM has a much smaller power density, thus temperature distribution on DRAM layers is primarily influenced by the underlying hot spots from the processor layer and the temperature variations are larger in bottom layers (hotter) than upper layers (cooler).

Furthermore, temperature drops as we move higher in the 3-D stack and closer to cooling. The temperature distribution shows large parts of a 3-D system will not exceed the temperature limitation during runtime. This observation provides an opportunity to improve the system performance by utilizing the




Fig. 3. Run-time temperature traces of CPU and DRAM with CPU-safe DTM and a trigger of 104°, running *gcc*.

biased temperature distribution in hot 3-D systems. We should note that the peak temperatures (not shown in the steady-state results), we observe during execution of *gcc* are 120 °C on the processor layer and 100 °C on the lowest DRAM layer. Thus, without any intervention, the processor and memory layers exceed their normal operating temperatures (by up to 15 °C).

B. Dynamic Thermal Management

We assume that our system uses a heat sink (which includes a fan) for cooling, located on top of the highest-stacked die. No other active cooling systems are deployed, such as porous silicon or liquid cooling [32]–[35], since they can be cost-prohibitive for 3-D implementations. When cooling solutions fail to promptly cool down the whole chip, modern processors employ DTM (e.g., DVFS) to address thermal emergencies. Typically, the system’s temperature is monitored periodically and once it exceeds a predefined threshold, DTM throttles the processor until temperature drops to safe levels. Due to the tight coupling between the CPU and memory parts in a 3-D system, DTM in a stacked system needs to consider the limitations of all layers as well as the interactions between them. In this paper, we assume the DTM throttles via conventional DVFS, dropping one DVFS step each time the threshold is reached.

1) *CPU Temperature Management*: We first quantify the problem of CPU-targeted DTM: we execute the *gcc* benchmark once again. We configure DTM (DVFS) to maintain the temperature across the processor layer below 105 °C (T_j max of several commercial processors [28]) and measure the *peak temperature* of our stacked DRAM during execution. Because temperature reacts slowly, DTM typically triggers at a temperature threshold slightly lower than its maximum operating temperature. In this motivating experiment, we set the trigger threshold at 104 °C. Fig. 3 presents the resulting temperature trace. We first observe that DTM successfully controls the processor layer temperature. Peak DRAM temperature, however, exceeds its operating temperature of 85 °C for a significant fraction of execution time and rises as high as 90 °C. During our experiment, we measure 12% of all memory accesses serviced by hot (unreliable) DRAM regions.

2) *DRAM Temperature Management*: As a result, DTM in a 3-D stacked system needs to consider the temperature requirements of the stacked DRAM. We expand the prior experiment using a similar DTM mechanism where the trigger

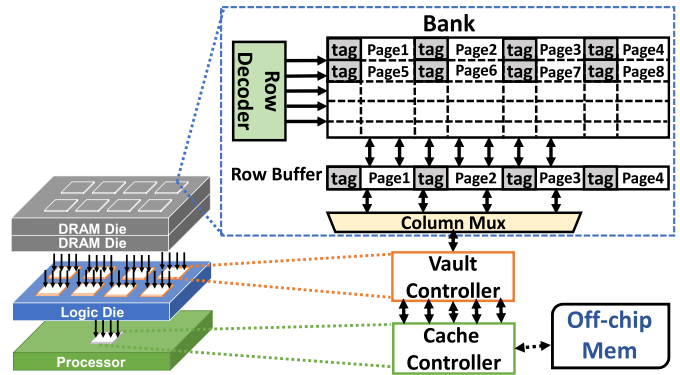


Fig. 4. 3-D-stacked DRAM cache (Unison).

temperature is set by DRAM temperature. Our experiments show that a trigger temperature set to 83 °C keeps all DRAM parts below 85 °C. However, DRAM-targeted DTM triggers much more often than the CPU-targeted DTM. Our experiments show the DRAM-targeted DTM causes an average 42% performance loss while CPU-targeted DTM only reduces the performance of the system by 12% on average (compared with the system without DTM). Therefore, traditional DTM, even when adjusted for DRAM, is not effective for 3-D systems.

III. TEMPERATURE-AWARE CACHE MANAGEMENT

By exploiting the observations in Section II, we can tolerate DRAM temperature violations while allowing the processor to run unthrottled, as long as we are not storing data in hot DRAM regions. We propose a temperature-aware mechanism which enables the system to work even when some parts of the DRAM are disabled.

A. DRAM Cache Organization

Our DRAM cache is organized as an HMC [1]. Based on HMC specifications, we use four DRAM and one logic layers, for a total LLC capacity of 4 GB. HMC is organized in 32 vaults that “slice” the four stacked memory chips vertically. A vault controller located on the logic layer manages traffic to and from all blocks, pages, and sets stored in its vault. We later exploit the independence of vault controllers to optimize our proposed mechanisms.

In this paper, we assume a page-based, tag-and-data (TAD) 3-D-stacked DRAM cache organization, managed using the Unison cache mechanism [12]. We implement a 4-way set associative LRU cache with four pages in each set. Page size is set at 2 KB, resulting in 8-KB cache sets, thus each set can be stored in a single DRAM row buffer, leading to improved access latency in the case of continuous accesses to the same set. Each page holds 31 data blocks; the one remaining block is used as a tag store. We further implement Unison’s block-based footprint miss predictor, allowing our cache to load only a subset of cache blocks, which are predicted to be accessed in the near future, inside a page when experiencing a miss. Fig. 4 illustrates the simplified model of our LLC organization.

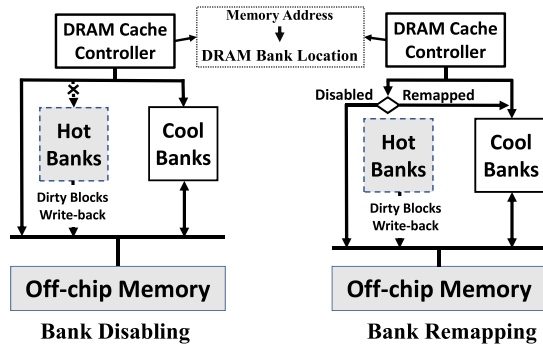


Fig. 5. Temperature-safe bank management.

B. Temperature-Safe Cache Operation

We introduce two temperature-safe methods to temporarily deactivate hot banks in DRAM cache: 1) DIS and 2) REM.

1) *Cache Disabling (DIS)*: Cache disabling is a straightforward solution to guarantee the absence of useful data in hot banks. DIS mechanism temporarily disables DRAM banks that exceed the thermal threshold. All accesses to cache blocks of the disabled banks are forwarded directly to the off-chip memory, as illustrated in Fig. 5(a). Since the disabled bank will no longer serve any memory operation, data retention is not an issue. Refresh can be entirely disabled in the bank to reduce power consumption without any reliability issue.

2) *Cache Remapping (REM)*: Disabling an entire (e.g., 16 MB in HMC [1]) bank will have a significant impact on cache miss rate, particularly if frequently accessed cache blocks become unavailable. Our remapping (REM) mechanism remaps cache blocks to cool DRAM areas that can reliably serve requests. While the cache still loses capacity, cache blocks can still find spots in the LLC.

The process of remapping is shown in Fig. 5(b). After remapping, the system will redirect the memory request of the original set to the target set, and look up the requested page by using the tag. In this situation, the number of pages the cache set might serve will be doubled. The cache needs to be distinguish the tags of two accesses originally mapped to different sets. As a result, the tag of each cache page is longer than that in a conventional cache not supporting remapping. Extra tag bits denote the original physical location in 3-D DRAM. New accesses to data that was previously in the hot bank will miss cache, loading the data into the new cool bank.

When one or both of the cache sets being combined are lightly accessed, the impact on cache miss rate should be minimal. If both are heavily utilized, we will see the impact of increased pressure on that combined cache page. Since it is possible that all banks are hot in the system, the REM mechanism also includes the *disabled* status to ensure all data are in temperature-safe places. Specifically, if a hot bank cannot find another bank to remap, it is disabled and accesses off-stack memory for the future data requests.

C. Exploration on Cache Remapping

Unlike the DIS mechanism, which is straightforward, there are several opportunities to optimize the REM policies.

1) *Target Selection*: A key policy decision is the selection of a target for the remapping. A straightforward target is the coolest bank in the system, which has a low probability of becoming hot in the future. This has two benefits. First, it minimizes the risk of remapping itself triggering a thermal event. Second, it reduces the likelihood of a complex remapping event where an oversubscribed bank is deactivated causing two banks' mappings to be adjusted.

Another potential target is a bank holding data which will have low access rate. The downside of bank-level remapping is reduced control over cache mappings in the target bank, however, it reduces hardware overheads for bookkeeping structures. At the bank level, we also observe that access rates are typically fairly even, particularly in the LLC after most of the locality has been hidden by the SRAM caches. So while it may be easy to find a block that is lightly used, it is difficult to find a bank of blocks that are all lightly used. Thus, we choose the coolest bank as the target when remapping. We further propose REM-G and REM-L which differ in the flexibility of bank target selection.

2) *Global Remapping (REM-G)*: In global remapping, the target for hot bank remapping is the coolest of all banks in the 3-D-stacked DRAM which was not previously used for remapping. This approach ensures that the optimal (temperature-wise) bank will be targeted for remapping each hot bank. To implement REM-G, a global remap table will be maintained in the DRAM cache controller to track the address remapping behavior for all banks.

3) *Local Remapping (REM-L)*: With REM-G, accessing the global remapping table is in the critical path of every memory request, which may degrade the performance even when there is no need to remap. With REM-L, we always choose a target bank in the same vault with the remapped bank. With this method, a local remap table is required in each vault controller on the 3-D-stacked DRAM logic layer. Each vault controller is in charge of address remapping for memory requests to this vault, which helps distribute and parallelize the cost of the table lookup. In addition, the local remap table (both individually and collectively) is smaller than the global remap table because of fewer bits required for the bank identifier. REM-L will suffer when an entire vault is much hotter than others, but that is unlikely because vaults span all layers of the DRAM stack, including those farthest from the CPUs.

4) *Remapping Level*: In this paper, we set a limit for level of remapping at two, which means at most two logical banks can be mapped to one physical bank. If the number of hot banks is over half of all banks, the new hot bank should be disabled. Furthermore, if a bank which has been combined with another bank becomes hot, both the banks need to be disabled. With remapping limited to two level, REM-G never ran out of banks, and REM-L reverted to DIS only rarely—so we do not consider higher levels of remapping.

D. Recovery

Since the temperature of any bank will vary based on CPU activities, especially when DTM is used to decrease system temperature, we must also be able to recover the usage of




Fig. 6. Average write-back latency over applications.

deactivated banks. For the DIS mechanism, it is quite straightforward. If we find that the temperature of a disabled bank is lower than a predefined recovery threshold, it will unmark the bank as disabled and enable the processor to start accessing cache sets in this bank again. This condition is also checked at each temperature sensing interval. We set the recovery threshold a little lower than the trigger threshold (85 °C) to avoid heavily vacillating between the two states.

For REM, if the temperature of one bank decreases under the recovery threshold, a recovery mechanism is invoked to enable cache sets again for the future access. Three situations could happen when a previous hot bank becomes cool.

- 1) If the bank has not been combined with other banks, the cache sets in the bank will be directly enabled.
- 2) If the bank is remapped to another bank, we will change the remap table and remap the bank back to its original physical location. In this situation, the dirty blocks that belong to the remapped sets need to be moved to the original locations.
- 3) If there was another bank remapped to the bank, we re-enable both of the original two bank mappings that use this bank.

E. Write-Back Optimization

In order to disable or remap a bank, all dirty cache blocks should be written back before any memory access to these data can be served. The process may lead to a significant performance loss if the number of dirty cache blocks becomes large. Fig. 6 shows the average write-back overhead when a global remapping (REM-G) is triggered in various workloads. Our results show the latency of write-back operations would range from tens of microseconds to over one millisecond. We explore an optimization mechanism to reduce the number of dirty cache blocks, thereby reducing the system-blocked time when triggering cache management.

Since temperature rises relatively slowly, we can typically identify potential hot pages early. We add an extra temperature threshold, which is slightly lower than the hot threshold (85 °C). Then, each bank in the system is in one of the three states based on two thresholds: 1) cool; 2) dangerous; and 3) hot. The main goal of managing data in dangerous banks is to reduce the significant overhead of write-back operations introduced by thermal-aware cache management mechanisms once the bank becomes hot. The simplest thing we can do is to treat dangerous banks as write-through. This will prevent adding to the number of dirty blocks, but not necessarily decrease it.

More proactively, we can write-back blocks on dangerous banks. Doing so immediately, however, just introduces the




Fig. 7. Comparison between original DTM and DTM with our proposed methods.

same overhead, stalling the cache or vault controllers. Instead, we leverage existing accesses to trigger write-back operations in dangerous banks. We can trigger data write-back exclusively on write accesses, or both read and write accesses. We can further write-back not just the cache block currently being accessed but also other cache data residing in the same DRAM row as the accessed block, such as a cache page or a cache set. We bound our write-back granularity to a DRAM row, since all its data are contained in a single (currently opened) DRAM row buffer. In this way, we avoid introducing additional overhead from issuing new row activation memory commands.

During the course of this paper, we explored various other combinations and options, spanning large ranges of three variables: 1) trigger temperature; 2) access trigger (only writes or both writes and reads); and 3) write-back granularities. We present the results of three different flavors of write-back optimizations in Section VI. *Writes-Block* essentially just turns the dangerous bank into a write-through cache. *Both-Page* and *Both-Set* trigger on both read and write accesses (row buffer reads), and vary in how aggressively they select blocks for write-back.

We should note none of the proposed operations on dangerous banks changes the already-presented thermal-aware cache management mechanisms. If a hot bank is recovered to serve cache accesses again and the temperature decreases to the “dangerous range,” we resume the write-back reduction policies. When the temperature decreases below the dangerous threshold, we can just resume the default cache operations.

IV. MECHANISM IMPLEMENTATION

A. Mechanism Overview

Our temperature-aware cache management methods are designed to complement DTM, not replace it. In fact, the two techniques are quite synergistic—our methods enable DTM to ignore temperature violations that will cause undue throttling of the CPU. In conventional DTM solutions, the system checks the temperature periodically and makes decision in each sensing interval. If the temperature exceeds the predefined trigger threshold, the system throttles to reduce its power consumption. The system recovers its performance when temperature drops below the safe threshold. Our mechanism is checked with the baseline DTM mechanism periodically with a same interval (one millisecond used in this paper).

As shown in Fig. 7, the trigger temperature is lower than the critical limit. As analyzed before, the temperature limit is usually the DRAM temperature limit and CPU is still safe at

this point. By adopting our methods in the existing DTM, the system can use a higher temperature threshold for throttling events while maintaining system reliability. Specifically, when the temperature reaches the original (DRAM-based) trigger point, the system needs to start protecting the DRAM. Instead of throttling the system, we deactivate the unreliable DRAM area. This allows system temperature to continue increasing to the new trigger temperature, which is set to protect CPU reliability. Only then will DTM kick in to cool down the system. For completeness, the figure shows the temperature decreasing to the lower DRAM trigger, at which point we restore the full capability of the DRAM cache. Overall, this approach allows the system to maintain high frequency over a much wider range of execution scenarios.

B. Temperature Sensing

Our proposed temperature-aware mechanisms operate based on a thermal representation of each memory die. For our exploration, we assume the existence of temperature sensors near each bank, which provides accurate thermal representation at the bank-level granularity. Our assumed implementation is not yet cost-effective, especially for 3-D-stacked architectures. However, researchers are already proposing solutions to increase the number of sensors [36], [37]. Prior works also propose models that can extrapolate accurate, fine-grained temperature sensing using the sensors already existing on the CPU layer in combination with the few sensors deployed on DRAM chips [38]–[40]. Finally, according to the HMC specification, there is unused space in the logic layer that could be used for temperature sensing. There are several recent works [41], [42] assuming a per-bank temperature sensing capability which utilizes the temperature sensors with small area overhead [43]. Furthermore, several temperature sensing techniques have been proposed to accurately sense temperature in 3-D ICs with low-cost [44]–[46]. Regardless of our assumption for ideal sensing capabilities, all our mechanisms can be easily adapted to operate based on extrapolated thermal models. In other words, the thermal representation is merely an input to our mechanism. The exact method of acquiring these measurements is not central to this paper.

C. Bookkeeping—Bank Control Table

Operating at bank-level granularity, our mechanisms require bookkeeping structures to keep track of state changes that lead to triggering of thermal-related actions. Particularly, we propose the addition of a “Bank Control Table.” With the exception of the *disable* and *dangerous* (single-bit) flags needed by all three mechanisms, bookkeeping requirements vary between our three proposed solutions. Furthermore, the placement of our BCT structure can affect performance; if placed right after the memory controller and before vault controllers it serializes accesses that could be completed in parallel. Centralized structures can also increase crossbar traffic and power consumption. Decentralized bookkeeping structures are often preferred in memory-related proposals. Fig. 8 illustrates the two possible locations to store our BCT.

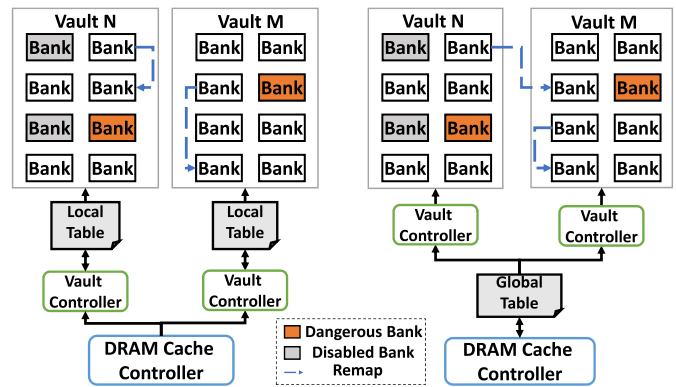


Fig. 8. Local (left) and global (right) bank control table.

Our simplest mechanism, DIS, only requires the two common flags (disable and dangerous), which are updated periodically on each sampling interval, depending on the choice of thermal thresholds set in our system. DIS requires minimal bookkeeping overhead: 512 bits are sufficient for our assumed 256 banks, distributed evenly across 32 vault controllers.

REM schemes require additional state information to keep track of remapped and merged banks. The *remap* field holds the physical bank location of this bank and is initialized to its bank’s ID. In case a bank is remapped, REM schemes update the field to reflect the change. When a remapped bank is accessed, the controller transparently redirects requests to the correct bank based on this field. We further add a *combine* flag, which uses a single bit to signal that some remote bank has been remapped to this bank. Using the combine flag, we reduce the complexity of remapping management by eliminating the need for reverse searches.

Local remapping (REM-L) requires three bits to identify the eight banks of a single vault in the remap field. Each vault controller maintains its own table since global bank-state information is unnecessary. REM-L requires a total of 224 bytes for BCT, evenly distributed across controllers (seven bytes each). Global remapping on the other hand requires 8-bit fields for bank IDs, leading to a total BCT overhead of 384 bytes. Compared to REM-L, REM-G incurs higher overhead for bookkeeping. Furthermore, a global BCT is conceptually centralized and must be accessed prior to any memory request. With REM-L, a portion of the remap table can be stored near each vault controller. The benefit of doing so is similar to the local DIS table, which can distribute and parallelize accesses to the table.

D. Controllers

To implement our proposed mechanism, a small amount of extra logic needs to be added in the DRAM cache controller and vault controllers. On an L2 cache miss, the on-chip memory controller sends a memory request to the DRAM cache controller to access the data. Since remap requires a longer page tag, the DRAM cache controller will generate a new tag combining the original tag and set number. The set number is used for the cache controller to find the data in the DRAM. The cache controller first extracts the bank and vault

number to get the status of the corresponding bank. Following operations depend on what mechanism the system uses.

Thermal-aware cache management operations related to each mechanism are checked every 1 ms, which is the temperature sensing and baseline DTM checking interval used in this paper. Based on the updated temperature, the corresponding controller (vault controller for local BCT or cache controller for global BCT) updates the BCT based on the mechanism and different temperature thresholds. For each cache access, if the *dangerous* flag of the target bank is set, the cache controller may initiate one or more write-backs depending on which write-back optimization is used.

E. Mechanism Processing Flow

1) *Thermal-Aware Cache Management*: Thermal-aware cache management operations are checked every 1 ms, which is the temperature sampling interval used in this paper. Based on the updated temperature, the corresponding controller updates the bank control table based on the mechanism and different temperature thresholds. In DIS, each vault controller identifies banks that are hot and not disabled. The vault controller notifies the cache controller to write-back all valid data in the bank and the vault controller sets the *disabled* flag. If a disabled bank becomes cool, the vault controller only needs to unset the disabled bit in the table.

In REM-G, if a bank is hot and it has neither been combined nor disabled, the cache controller selects the coolest available bank and updates the remap table. If REM-G fails to find a target bank, the disabled bit in the global REM table will be set by the cache controller, which then writes back dirty cache blocks in the hot bank. The remapping information will not be changed in that case. At the same time, if one disabled bank becomes cool again, the cache controller would change the remapping information in the REM table based on the recovery mechanism. REM-L is similar to REM-G in this scenario except that each vault controller handles these operations.

2) *Cache Access*: In DIS, each vault controller can access a local (per-vault) DIS table to control the mechanism. Once a memory request comes to a vault controller, it will first check whether the corresponding bank is disabled. If it is disabled, the cache controller directly accesses the data in the off-chip memory. In REM-G, the DRAM cache controller needs to access the global REM table on every memory request, and check whether the corresponding bank is disabled or remapped to another bank. If the bank is disabled, we directly access off-chip memory for the requested data. If the bank is remapped to another bank, the cache controller changes the address bits which indicate the bank index. It then sends the request to the corresponding vault controller of the remapped bank. REM-L is similar to REM-G except all operations are processed by vault controllers.

3) *Write-Back on Access*: For each cache access, if the *dangerous* flag of the target bank is set, the cache controller may initiate one or more write-backs. If the cache access causes a cache miss, the cache controller would only issue extra commands when the set-level write-back is used. If there is a cache

hit, the cache controller would write-back dirty cache blocks before the activated DRAM row is replaced.

4) *Protection for Hot Banks*: To protect the data in the hot bank during the disabling and remapping process, we set the trigger temperature lower than the critical DRAM temperature limit. We empirically choose the trigger temperature such that the temperature of a hot bank will not reach the critical temperature limit before finishing the disabling or remapping operation.

F. Discussion on DRAM Cache Technologies

Besides the Unison cache, more DRAM cache managing techniques can be found in the literature such as Alloy cache [13]. Our proposed mechanisms are largely orthogonal to the choice of DRAM cache design. Our sole requirement is the existence of translation from memory address to DRAM location (bank-level). For different DRAM cache organization, the cache controller still needs to extract a DRAM vault and bank address from each memory request. Thus, the BCT should be the same. Since we store the full tag for each cache block, the remapping mechanism will also be effective if we only change the address bits indicating the physical DRAM location. For write-back optimization, different cache organizations will support different granularities of write-back on access.

Furthermore, the several DRAM caches, including Unison cache, deploy a prefetching mechanism to reduce the overhead of cache miss [12], [47]. In our design, prefetching techniques will remain operational, even during the short temporary increase of memory traffic caused when our mechanisms react to a thermal emergency, and may in fact help hide those delays. The coexistence of prefetchers alongside our mechanisms does not require additional management logic.

V. METHODOLOGY

A. Simulation Infrastructure

We address the experimental needs of this paper using a combination of Sniper [48], Ramulator [49], Cacti [50], McPat [51], and Hotspot [52]. During our experiments, control is transferred between these five tools, based on a series of carefully synchronized, event-based software interrupts, which form the core of our modifications.

Sniper is a multicore simulator based on the interval core model and the graphite simulation infrastructure [53]. We extend Sniper and integrate it with Ramulator: a cycle-accurate DRAM simulator. We configure Ramulator according to our stacked DRAM cache and off-chip memory configurations and we extend it to handle the stacked portion of DRAM as Unison cache. The Sniper–Ramulator combination is sufficient for our performance measurements. We interrupt Sniper on each L2 cache miss, at which point control is transferred to the Ramulator module. Following this control transfer, the internal LLC controller performs necessary transformations dependent on our bookkeeping structures, and finally, services the request as a typical DRAM controller. Table I presents the configuration of systems tested in our experiments. In the 3-D setting, the aggregated bandwidth provided by 32-vault channels is

TABLE I
SYSTEM CONFIGURATION

	Configuration
Technology	22nm, 1.0V, 2.66GHz
Cores	8 @ 2.66GHz Intel Nehalem-like
Cache	32 KB L1 Cache, 256 KB L2 Cache
Off-chip DRAM	4 channels, 7.6GB/s per channels 4 ranks per channel, 8 banks per rank
Stacked DRAM	4GB, 500MHz bus frequency 32 channels, 16MB banks, 8KB row buffer internal: 512GB/s, external: 320GB/s (2Gb/s TSVs)
t_{CAS} - t_{RCD} - t_{RP} - t_{RAS}	7-7-7-17
t_{RC} - t_{WR} - t_{WTR} - t_{RTP}	24-8-4-7
t_{RRD} - t_{FAW}	5-20

TABLE II
THERMAL PARAMETERS

Thermal interface material	
Specific Heat Capacity	$4 \times 10^6 \text{ J/m}^3\text{K}$
Resistivity	0.25 mK/W
Thickness	0.02 mm
TSV	
Specific Heat Capacity	$4 \times 10^6 \text{ J/m}^3\text{K}$
Resistivity	0.0058 mK/W
Thickness	0.02 mm
Processor and DRAM Silicon	
Specific Heat Capacity	$1.75 \times 10^6 \text{ J/m}^3\text{K}$
Resistivity	0.01 mK/W
Thickness	0.15 mm
Heat Sink Specification	
Convection Capacitance	140.4 J/K
Convection Resistance	0.1 K/W
Thermal Conductivity	400 W/mK

512 GB/s. In the 2.5-D setting, the HMC chip is connected with the host processor by four 16-lane, full-duplex serialized links, which provide up to 320-GB/s bandwidth. This data comes from the official HMC specification [1].

We incorporate McPat and Cacti for all our power consumption measurements. We use Cacti to estimate area overhead and power consumption of our SRAM-based bookkeeping structures. Similarly, we use McPat for power estimation for logic dies (processor and DRAM logic layer) using 22-nm technology. The DRAM power consumption is obtained by the published industry DDR4 specification sheets [54]. We model four vertically stacked DRAM layers and one processor layer with eight cores. Prior to running any experiment, we calibrated our McPat configuration such that results match those of prior work [55].

We utilize HotSpot to acquire runtime temperature traces of the system. HotSpot periodically interrupts simulations and calculates the transient temperature based on prior thermal traces and the new interval's power consumption measurements. We configure HotSpot intervals to 1 ms of simulated time. Table II presents our HotSpot configuration in detail, including thermal interface material (layers of material residing between stacked dies for heat dissipation), TSV thermal characteristics, and the assumed cooling system. The material characteristics of thermal interface material, TSV, DRAM, and silicon are from a validated previous work [56].

B. Workloads and Baseline Processor

For our simulations, we use ten-billion instruction-long benchmark phases from the SPEC2006 benchmark suite [57].




Fig. 9. DTM Performance and hot-bank access rate for two sensor targets (CPU or memory) and a variety of temperature triggers.

Most SPEC2006 benchmarks have very small memory footprints, making it challenging to stress-test a 4-GB DRAM cache. We address this issue following two approaches: first, we simulate eight copies of the same benchmark running simultaneously on our 8-core system, essentially multiplying its memory footprint eightfold. We find that various benchmarks exceed our LLC capacity using this model. Second, we present cache size sensitivity experiments, going as far as eliminating the stacked LLC entirely, providing a fair baseline comparison. Our processor layer configuration is presented in Table I.

C. Baseline DTM

The DVFS technology used in this paper has five voltage-frequency settings. Frequency decreases by 20% when the V/f setting is decreased by one step. The minimum voltage is 70% of the maximum one which is based on published data of commercial processors [58]. When the peak temperature of one layer exceeds the trigger temperature, the system will decrease the V/f setting by one step unless it is already the lowest one. In our experiments, the lowest V/f is enough to decrease the temperature of the system without further throttling. When the temperature drops below the trigger temperature, the system increases the V/f setting by one step. Since temperature of both processor and DRAM can affect DTM, we distinguish different DTM policies based on which temperature threshold is used. For example, CPU-105 means the DTM triggers when any part of the processor becomes hotter than 105 °C. At the same time, DRAM-85 means DTM triggers when any part of DRAM (not including the DRAM logic layer) becomes hotter than 85 °C. In simulation, the latency of changing V/f setting is assumed to be instantaneous—this favors the DTM baseline, which changes frequency more frequently.

VI. EXPERIMENTAL RESULTS

A. Baseline DTM

To understand the effectiveness and cost of a conventional DTM, we examine a sweep of DTM settings, all using the same basic algorithm described in Section V. The variables are: 1) whether we are using CPU temperature or DRAM temperature as our trigger and 2) the trigger temperature. Fig. 9 shows the result of this exploration, averaged across all applications. The performance result is shown as an average IPC of eight cores in the system, normalized to average IPC of the system without any thermal reactive mechanism. We should




Fig. 10. Overall performance of different DTM and cache mapping mechanisms.



Fig. 11. Overall power consumption of different DTM and cache mapping mechanisms.

note that the IPC results are all scaled to a common cycle time at the end of simulation, which means the IPC results shown can be used to compare results across runs with different DVFS behavior. Hot access rate indicates what percentage of memory requests access data stored in a DRAM area whose temperature is larger than 85 °C. For DRAM-targeted DTM (where we monitor DRAM peak temperature, and use that to drive DTM), the range of trigger temperatures is from 80 °C to 90 °C. The result shows the highest trigger temperature for DRAM-target DTM to avoid any hot access in any application is 83 °C. At 84 °C, the hot access rate is low, which is about 0.1%, but not zero. However, the DRAM-Safe (DRAM83) mechanism suffers 37.2% performance loss.

For CPU-targeted DTM, the tested range of trigger temperatures is from 95 °C to 105 °C. In this paper, the T_j max of the processor is assumed to be 105 °C. We can drive the trigger temperature pretty close to T_j max without violating CPU limits. In that case, however, many DRAM accesses are unsafe (12.6% and 11.7% in CPU-105 and CPU-104, respectively). We can set the CPU trigger temperature lower, but even at 95 °C and 96 °C, the unsafe DRAM accesses are 0.1% and 0.8% while sacrificing 25.7% and 24.0% performance, respectively. We would have to set the CPU trigger lower still to completely eliminate unsafe DRAM accesses, incurring greater performance loss. Based on our experiment, setting CPU trigger to 88 °C can eliminate DRAM temperature violations entirely, but incurs a performance loss of 42%.

B. Overall Results

Based on the results of the baseline DTM, we use DRAM83 as a comparison point and relabel it as our *DRAM-Safe* baseline. We will also label CPU104 as CPU-Safe, but we must keep in mind that CPU-Safe is safe for the CPU but not for memory. It serves as an upper bound for reasonable

performance. We will also consider IDEAL, which is the full performance with no DTM throttling whatsoever—also unsafe, and a higher upper bound. For more comprehensive comparison, a no-LLC baseline and 2.5-D baseline are added in this experiment. Specifically, no-LLC baseline removes 3-D-stacked DRAM cache from the tested system and any L2 cache miss will access off-chip memory. The 2.5-D baseline still utilizes stacked DRAM as LLC but puts it off-chip with a corresponding higher bus delay. Both no-LLC and 2.5-D baseline do not cause DTM to ever be invoked in our simulations.

We will combine our three DRAM management techniques with CPU-Safe, attempting to eliminate the unsafe accesses to DRAM of CPU-Safe without the severe performance loss of DRAM-Safe. Fig. 10 shows the performance result of the eight configurations on all applications. Based on the results, all of the three proposed mechanisms outperform the DRAM-Safe approach. Specifically, DIS, REM-L, and REM-G show an average performance improvement over DRAM-Safe of 21.3%, 23.4%, and 26.1%, respectively. Our mechanisms introduce only a small overhead relative to the CPU-Safe DTM upper bound, being within 8.1%, 6.6%, and 4.5%, respectively. Based on the results, our proposed cache management methods can still outperform two no-DTM baseline system. REM-G achieves 11.3% and 22.9% average performance improvement over the 2.5-D baseline and the no-LLC baseline, respectively. The latter result was run to verify that our gains are not coming simply because our workloads make poor use of the LLC.

Fig. 11 shows the power consumption of different mechanisms used in this paper. Compared to the ideal solution without any DTM, the normalized power consumption of DIS, REM-L, and REM-G are 0.9, 0.94, and 0.97, respectively. We also calculate the energy consumption based on the power and execution time, our proposed mechanisms consume 7.1%, 9.8%, and 11.0% more energy than the ideal architecture. Such energy consumption overhead mainly comes from




Fig. 12. AMAL normalized to CPU-Safe baseline.




Fig. 13. Cache miss results of different mechanisms.

the extra memory operations introduced by our mechanisms. As a comparison, DRAM-safe mechanism consumes 9.5% more energy than the ideal architecture. Such results show that our mechanisms provide a comparable energy efficiency with DRAM-safe mechanism, while significantly increasing the performance. The results also show that these mechanisms control the temperature by effectively limiting the power consumption of the whole system. The extra power is consumed by the memory chip, which contributes much less to the temperature than CPUs.

C. Memory Access Latency

We further investigate the average memory access latency (AMAL) of three methods and compare them with that of the CPU-Safe baseline. Fig. 12 shows the normalized AMAL for each method. The average increase in AMAL of DIS, REM-L, and REM-G are $2.17\times$, $1.86\times$, and $1.50\times$, respectively. To explain such AMAL increase, the cache miss rate results are shown in Fig. 13. The geometric means of cache miss rate over all workloads are 21.4%, 17.6%, and 11.2% for DIS, REM-L, and REM-G, respectively. Furthermore, Fig. 14 shows the frequency of management events for all methods. We record the number of block misses, page misses, and accesses to disabled pages. The results show DIS causes the most block misses because we disable all hot banks. REM-G never disables a bank, which means it can always find a cool bank to remap to. However, remapping cache sets leads to a significant increase in page misses. Since Unison cache employs a footprint predictor to load partial data of a page during a page miss, the negative effect of the page miss is reduced. These results show the three methods have different memory behaviors because of different management policies. Overall, REM-G outperforms the other two methods primarily because of lower cache miss rate resulting from increased flexibility in remapping. The reason why REM-G can always find a target bank for remapping is that the temperature distribution is uneven in not only the vertical but also the horizontal direction. In that case, REM-G remaps the hot banks to the cool banks unless the number of the hot banks in the memory system is larger than that of the cool banks.





Fig. 14. Frequency of cache management operations of different mechanisms.

Fig. 15. Results of IPC and miss rate with various CPU-target DTM trigger temperature in *gcc* (left) and *bzip2* (right). For each threshold, we test four configurations (from left to right): DIS, REM-L, REM-G, and CPU-Safe.

D. Temperature-Sensitivity Experiment

The choice of trigger temperature is a less obvious one, when combined with our proposed cache management methods, compared to the conventional case. A higher trigger temperature will minimize throttling events in the CPU, but also maximize the number of DRAM banks that are deactivated. More deactivated DRAM banks will minimize LLC capacity and LLC hit rate. Lower trigger temperatures, then, will increase throttling but reduce average memory access time.

Fig. 15 shows the results of temperature sensitivity experiments on a memory-moderate workload (*gcc*) and a memory-intensive workload (*bzip2*). Both IPC and cache miss rate results for REM-G are better than REM-L and DIS. REM-L also has a higher IPC and lower cache miss rate than DIS. However, in decreasing the trigger temperature for DTM, the gaps for IPC and cache miss rate between different methods shrink because there are fewer banks being deactivated due to the CPU's lower overall temperature. Overall, though, these results confirm that despite a drop in LLC miss rate, we achieve the highest throughput with aggressive DTM temperature triggers, sacrificing some memory access latency in exchange for unfettered CPU performance.

E. Write-Back Optimization

The previous results show that the three proposed mechanisms have already improved the performance over the baseline DRAM-Safe system significantly. However, our results indicate we are still losing performance due to high write-back activity, so there is still an opportunity. Because REM-G shows the best overall performance based on the previous experiments, we use it as the baseline method in this section. We now introduce the notion of a *dangerous* page, as described in Section III-E.




Fig. 16. Average write-back latency reduction during each cache management in REM-G.

We test three dangerous temperature thresholds: 80 °C, 82 °C, and 84 °C. For each temperature threshold, we test three mechanisms: 1) *Write-Block*; 2) *Both-Page*; and 3) *Both-Set*. Write-block essentially turns the page into a write-through cache for the future accesses. The other two potentially trigger multiple write-backs after any LLC access. Fig. 16 shows the average write-back latency reduction. Because of the space limitation, we show results for four different workloads which show different patterns of behavior, each representative of a subset of the full suite. All mechanisms only eliminate a small portion of the latency in *bzip2*, but reduce over 40% in *gcc*. Furthermore, the latency reduction of *milc* is sensitive to access type while *bwaves* is sensitive to temperature threshold.

We then look into the overall IPC improvement of write-back mechanisms. Fig. 17 shows the best performance improvement provided by the tested nine mechanisms, which is *T80-Both-Page*, compared with REM-G, as well as ideal REM-G which assumes all write-back latencies are zero. On average, we see that the additional gains due to write-back mitigation are not insignificant, but not large; however, in general we achieve half to a third of the upper bound available gains. We found that choosing the best of our policies for each benchmark actually does quite a bit better than this, but that implies a more complex controller.

F. Cache Size Sensitivity Experiments

The pressure on the LLC varies by workload and is also impacted by our choice of an eight-core CPU for our experiments. However, increased cache pressure (which can change the tradeoff between cache size reduction and CPU throttling that is key to this paper) can come from several sources, including more intensive workloads, more cores, more threads per core, etc. To capture the effects of an LLC under greater pressure, we examine a variety of LLC sizes.

Specifically, we scale the size of each bank, while keeping the basic HMC structure which consists of 32 vaults and each vault having eight banks, to provide maximum consistency with our other results. Fig. 18 shows the IPC and cache miss rate result of one compute-intensive workload (*gcc*) and one memory-intensive workload (*bzip2*) on different configurations—other results are similar. For the compute-intensive workload, small cache size significantly increases the LLC miss rate and therefore decrease the IPC for all configurations; however, the incremental loss in cache miss rate due to REM-G stays fairly consistent even with small caches. Because we are only impacting the LLC and not the L1 and




Fig. 17. Overall performance improvement of our most effective write-back mitigation mechanism, *T80-Both-Page*.




Fig. 18. Result of IPC and cache miss rate of different configurations with different cache size over *gcc* (compute intensive) and *bzip2* (memory intensive.)

L2 caches, the impact of the higher LLC miss rates is still somewhat muted. Overall, this means that we continue to see the effectiveness of REM-G even in the presence of higher LLC pressure.

G. Comparison With Thermal-Aware Refresh Management

Increasing the refresh rate of DRAM is the traditional method to protect DRAM from the faster data loss caused by high temperature. However, increasing the frequency of refresh operations introduces large overhead in terms of performance and energy consumption in DRAM [23], [59]. Moreover, the negative impact caused by refresh operations may become more severe if the capacity of DRAM increases. Finally, there is no existing technique in the literature that applies different refresh rates in different DRAM areas experiencing different thermal situations. Thus, the refresh rate of the whole DRAM chip would be increased if any part of the memory becomes hot. We then compare our proposed mechanisms with the system which doubles the refresh rate when the temperature exceeds 85 °C.

The energy overhead of increasing refresh rate is estimated based on the execution time of the high temperature phase and the energy consumption of a refresh operation based on published data [54], [59]. For these experiments, we estimate the memory access latency overhead caused by increasing refresh rate after each simulation interval. The estimation is based on the temperature and memory access patterns during the last interval and the expected rate of conflict between access and refresh. The frequency of checking the temperature is the same, and the energy consumption includes the energy consumed by both the logic and the DRAM layers. Fig. 19 shows the performance and memory energy consumption of increasing refresh rate compared to REM-G with *T80-Both-Page* write-back optimization. Because of more frequent refresh operations, thermal-aware DRAM refresh mechanism consumes 7.4% more memory energy in average over all tested




Fig. 19. Comparison between REM-G (with T80-Both-Page optimization) and thermal-aware DRAM refresh mechanism. The left (blue) bar and right (red) bar indicate the energy consumption and the IPC of REM-G, respectively. All values are normalized to corresponding values of the thermal-aware DRAM refresh mechanism.

applications. Furthermore, enabling a double refresh rate does not provide a better performance than our proposed mechanisms. The results show that our best mechanism provides a slight (1.5%) speedup over the thermal-aware DRAM refresh mechanism while saving a considerable amount of energy.

H. Discussion About Multithreaded Workloads

This paper focuses on multiprogram workloads, and analysis of multithreaded (parallel) workloads is left for the future work. However, we expect these solutions will have similar or even better performance when running a multithreaded workload. In particular, in this paper, we have modeled a homogeneous multiprogrammed workload. We do this because it stresses our experiments in two ways: 1) it tends to heat each core similarly, which reduces our opportunity to remap (a more diverse workload, and resulting thermal pattern, would allow us even more opportunity to remap successfully) and 2) it maximizes pressure on the LLC (our solutions achieve highest gain when LLC pressure is low). With a parallel workload, we would tend to see the same uniform thermal pattern we see in these experiments, but typically lower LLC pressure due to the nonlinear scaling of data size on parallel executions [60].

VII. RELATED WORK

Several studies in the literature focus on thermal issues in 3-D stacked systems. Eckert *et al.* [61] explored the thermal feasibility of processing in memory (PIM) based on 3-D-stacked DRAM, and show that PIM is thermally feasible even with low-end fanless cooling solutions. However, that work focuses on processors with very low power consumption, while this paper explores the impact of 3-D thermals in the context of state-of-the-art high-performance processors. Zhu *et al.* [62] investigated the thermal issues existing in processing the die-stacking memory and showed that the host CPU dominates the thermals of a system with a die-stacked PIM accelerator. Unlike their work, this paper assumes a DRAM stack placed directly on the top of the host CPU, which introduces more critical thermal challenges in die-stacking systems. Coskun *et al.* [32] utilized microchannel-based liquid cooling technology to control the thermal issues in a 3-D stacked architecture and propose a controller to adjust the liquid flow rate to minimize pump energy consumption. Their work is therefore orthogonal to conventional DTM mechanisms, which likely still need to be deployed.

Several works have examined temperature-aware management in memory systems [21], [34], [56], [63]–[68]. Liu *et al.* [21] proposed three hardware and software schemes to reduce peak temperature on a traditional DRAM chip. Kang *et al.* [63] applied runtime cache tuning with per-core DVFS to maximize the performance of chip multiprocessors with 3-D-stacked LLC memory without thermal-constraint violation. Meng *et al.* [56] proposed a runtime optimization policy to maximize performance while maintaining power and thermal constraints. TAPAS [64] is a low-cost temperature-aware adaptive block placement and migration policy to reduce access to hot banks for hybrid LLC consisting of STT-RAM and STT-SRAM. There are also several papers focusing on controlling the temperature in systems with die-stacking memory using pipeline control [68], thread migration [65], page allocation [66], and adaptive DVFS [67]. All these previous works focus on peak temperature reduction and cannot completely remove the DTM in the current system. On the contrary, our proposed mechanisms are triggered upon thermal emergencies to improve the performance of conventional DTM and guarantee that DRAM serves memory requests reliably. Furthermore, our methods can be triggered when these temperature-reduction proposals fail to limit the temperature.

Furthermore, disabling parts of the cache was first introduced to reduce cache energy consumption by Albonesi [69]; but the motivation and the cache organization are completely different, necessitating quite different approaches.

VIII. CONCLUSION


This paper explored thermal issues in a 3-D system with stacked-DRAM as a large LLC combined with a CPU logic layer. This configuration maximizes the throughput and effectiveness of the stacked DRAM for high-performance systems, but presents thermal challenges—not just from the cooling standpoint but also due to the mix of temperature sensitivities between CPU and DRAM. The thermal coupling between the layers, and the different temperature limits for CPU and DRAM make conventional DTM ineffective for controlling temperature in 3-D systems. Thus, we propose three DRAM cache management mechanisms, which work with conventional DTM. These allow DTM to only throttle the system for CPU thermal events, instead sacrificing LLC space to eliminate DRAM-induced CPU throttling, enabling the system to maintain close to full performance even in the face of high utilization and high thermal activity. The results show that the proposed mechanisms can improve the performance over DRAM-Safe DTM by up to 26.1% on average which makes 3-D systems more practical for the future high-performance computing.

REFERENCES

- [1] *Hybrid Memory Cube Specification 2.1*. Accessed: Nov. 2018. [Online]. Available: <http://hybridmemorycube.org/specification-v2-download/>
- [2] *High Bandwidth Memory (HBM) Dram*. Accessed: Oct. 2018. [Online]. Available: <https://www.jedec.org/standards-documents/docs/jesd235a>
- [3] E. Solomonik, M. Besta, F. Vella, and T. Hoefler, “Scaling betweenness centrality using communication-efficient sparse matrix multiplication,” in *Proc. ACM Int. Conf. High Perform. Comput. Netw. Storage Anal.*, Nov. 2017, Art. no. 47.

- [4] T. Kurth *et al.*, “Deep learning at 15PF: Supervised and semi-supervised classification for scientific data,” in *Proc. ACM Int. Conf. High Perform. Comput. Netw. Stor. Anal.*, 2017, p. 7.
- [5] A. Li *et al.*, “Exploring and analyzing the real impact of modern on-package memory on HPC scientific kernels,” in *Proc. ACM Int. Conf. High Perform. Comput. Netw. Stor. Anal.*, Denver, CO, USA, 2017, p. 26.
- [6] J. Ouyang *et al.*, “Arithmetic unit design using 180nm TSV-based 3D stacking technology,” in *Proc. IEEE Int. Conf. 3D Syst. Integr. (3DIC)*, San Francisco, CA, USA, 2009, pp. 1–4.
- [7] J. Zhao, Q. Zou, and Y. Xie, “Overview of 3-D architecture design opportunities and techniques,” *IEEE Des. Test.*, vol. 34, no. 4, pp. 60–68, Aug. 2017.
- [8] *An Intro to MCDRAM (High Bandwidth Memory) on Knights Landing*. Accessed: Jan. 2016. [Online]. Available: <https://software.intel.com/en-us/blogs/2016/01/20/an-intro-to-mcdram-high-bandwidth-memory-on-knights-landing>
- [9] *Nvidia Tesla v100—Nvidia*. Accessed: Jan. 2019. [Online]. Available: <https://www.nvidia.com/en-us/data-center/tesla-v100/>
- [10] *AMD Firepro S9300 X2 Server GPU*. Accessed: Oct. 2018. [Online]. Available: <https://www.amd.com/en-us/products/graphics/server/s9300-x2>
- [11] J. Jeddeloh and B. Keeth, “Hybrid memory cube new DRAM architecture increases density and performance,” in *Proc. IEEE Symp. VLSI Technol. (VLSIT)*, Honolulu, HI, USA, 2012, pp. 87–88.
- [12] D. Jevdjic, G. H. Loh, C. Kaynak, and B. Falsafi, “Unison cache: A scalable and effective die-stacked DRAM cache,” in *Proc. 47th Annu. IEEE/ACM Int. Symp. Microarchit.*, Cambridge, U.K., 2014, pp. 25–37.
- [13] M. K. Qureshi and G. H. Loh, “Fundamental latency trade-off in architecting DRAM caches: Outperforming impractical SRAM-tags with a simple and practical design,” in *Proc. 45th Annu. IEEE/ACM Int. Symp. Microarchit.*, Vancouver, BC, Canada, 2012, pp. 235–246.
- [14] Y. Lee *et al.*, “A fully associative, tagless DRAM cache,” *ACM SIGARCH Comput. Archit. News*, vol. 43, no. 3, pp. 211–222, 2015.
- [15] X. Yu, C. J. Hughes, N. Satish, O. Mutlu, and S. Devadas, “Banshee: Bandwidth-efficient DRAM caching via software/hardware cooperation,” in *Proc. 50th Annu. IEEE/ACM Int. Symp. Microarchitect.*, Cambridge, MA, USA, Oct. 2017.
- [16] C. C. Chou, A. Jaleel, and M. K. Qureshi, “CAMEO: A two-level memory organization with capacity of main memory and flexibility of hardware-managed cache,” in *Proc. 47th Annu. IEEE/ACM Int. Symp. Microarchit.*, Cambridge, U.K., 2014, pp. 1–12.
- [17] M. R. Meswani, S. Blagodurov, D. Roberts, J. Slice, M. Ignatowski, and G. H. Loh, “Heterogeneous memory architectures: A HW/SW approach for mixing die-stacked and off-package memories,” in *Proc. IEEE 21st Int. Symp. High Perform. Comput. Archit. (HPCA)*, Burlingame, CA, USA, 2015, pp. 126–136.
- [18] A. Prodomou, M. Meswani, N. Jayasena, G. Loh, and D. M. Tullsen, “MemPod: A clustered architecture for efficient and scalable migration in flat address space multi-level memories,” in *Proc. IEEE Int. Symp. High Perform. Comput. Archit. (HPCA)*, Austin, TX, USA, 2017, pp. 433–444.
- [19] T. Hamamoto, S. Sugiura, and S. Sawada, “On the retention time distribution of dynamic random access memory (DRAM),” *IEEE Trans. Electron Devices*, vol. 45, no. 6, pp. 1300–1309, Jun. 1998.
- [20] J. Liu, B. Jaiyen, Y. Kim, C. Wilkerson, and O. Mutlu, “An experimental study of data retention behavior in modern DRAM devices: Implications for retention time profiling mechanisms,” *ACM SIGARCH Comput. Archit. News*, vol. 41, no. 3, pp. 60–71, 2013.
- [21] S. Liu, B. Leung, A. Neckar, S. O. Memik, G. Memik, and N. Hardavellas, “Hardware/software techniques for DRAM thermal management,” in *Proc. IEEE 17th Int. Symp. High Perform. Comput. Archit. (HPCA)*, San Antonio, TX, USA, 2011, pp. 515–525.
- [22] J. Mukundan, H. Hunter, K.-H. Kim, J. Stuecheli, and J. F. Martínez, “Understanding and mitigating refresh overheads in high-density DDR4 DRAM systems,” *ACM SIGARCH Comput. Archit. News*, vol. 41, no. 3, pp. 48–59, 2013.
- [23] J. Liu, B. Jaiyen, R. Veras, and O. Mutlu, “RAIDR: Retention-aware intelligent DRAM refresh,” *ACM SIGARCH Comput. Archit. News*, vol. 40, no. 3, pp. 1–12, 2012.
- [24] J. Lim, H. Lim, and S. Kang, “3-D stacked DRAM refresh management with guaranteed data reliability,” *IEEE Trans. Comput.-Aided Design Integr. Circuits Syst.*, vol. 34, no. 9, pp. 1455–1466, Sep. 2015.
- [25] Z. Liu, T. Xu, S. X.-D. Tan, and H. Wang, “Dynamic thermal management for multi-core microprocessors considering transient thermal effects,” in *Proc. IEEE 18th Asia South Pac. Design Autom. Conf. (ASP-DAC)*, Yokohama, Japan, 2013, pp. 473–478.
- [26] J. Srinivasan and S. V. Adve, “Predictive dynamic thermal management for multimedia applications,” in *Proc. ACM 17th Annu. Int. Conf. Supercomput.*, San Francisco, CA, USA, 2003, pp. 109–120.
- [27] D. Brooks and M. Martonosi, “Dynamic thermal management for high-performance microprocessors,” in *Proc. 7th IEEE Int. Symp. High Perform. Comput. Archit. (HPCA)*, 2001, pp. 171–182.
- [28] *Intel. Intel Core i7–3612QM Processor*. Accessed: Oct. 2018. [Online]. Available: http://ark.intel.com/products/67356/Intel-Core-i7-3612QM-Processor-6M-Cache-up-to-3_10-GHz-rPGA
- [29] B. Schroeder, E. Pinheiro, and W.-D. Weber, “DRAM errors in the wild: A large-scale field study,” *ACM SIGMETRICS Perform. Eval. Rev.*, vol. 37, no. 1, pp. 193–204, 2009.
- [30] A. Agrawal, J. Torrellas, and S. Idgunji, “Xylem: Enhancing vertical thermal conduction in 3D processor-memory stacks,” in *Proc. 50th Annu. IEEE/ACM Int. Symp. Microarchit.*, Boston, MA, USA, 2017, pp. 546–559.
- [31] J. Long, S. O. Memik, G. Memik, and R. Mukherjee, “Thermal monitoring mechanisms for chip multiprocessors,” *ACM Trans. Archit. Code Optim.*, vol. 5, no. 2, p. 9, 2008.
- [32] A. K. Coskun, D. Atienza, T. S. Rosing, T. Brunschweiler, and B. Michel, “Energy-efficient variable-flow liquid cooling in 3D stacked architectures,” in *Proc. Conf. Design Autom. Test Europe*, Dresden, Germany, 2010, pp. 111–116.
- [33] S. G. Kandlikar, “Review and projections of integrated cooling systems for three-dimensional integrated circuits,” *J. Electron. Packag.*, vol. 136, no. 2, 2014, Art. no. 024001.
- [34] M. V. Beigi and G. Memik, “TESLA: Using microfluidics to thermally stabilize 3D stacked STT-RAM caches,” in *Proc. IEEE 34th Int. Conf. Comput. Design (ICCD)*, Scottsdale, AZ, USA, 2016, pp. 344–347.
- [35] Y. Zhang, L. Zheng, and M. S. Bakir, “3-D stacked tier-specific microfluidic cooling for heterogeneous 3-D ICS,” *IEEE Trans. Compon. Packag. Manuf. Technol.*, vol. 3, no. 11, pp. 1811–1819, Nov. 2013.
- [36] D. Shim, H. Jeong, H. Lee, C. Rhee, D. K. Jeong, and S. Kim, “A process-variation-tolerant on-chip CMOS thermometer for auto temperature compensated self-refresh of low-power mobile DRAM,” *IEEE J. Solid-State Circuits*, vol. 48, no. 10, pp. 2550–2557, 2013.
- [37] C.-K. Kim, J.-G. Lee, Y.-H. Jun, C.-G. Lee, and B.-S. Kong, “CMOS temperature sensor with ring oscillator for mobile DRAM self-refresh control,” *Microelectron. J.*, vol. 38, no. 10, pp. 1042–1049, 2007.
- [38] S. Sharifi and T. Š. Rosing, “Accurate direct and indirect on-chip temperature sensing for efficient dynamic thermal management,” *IEEE Trans. Comput.-Aided Design Integr. Circuits Syst.*, vol. 29, no. 10, pp. 1586–1599, Oct. 2010.
- [39] Y. Zhang and A. Srivastava, “Accurate temperature estimation using noisy thermal sensors for Gaussian and non-Gaussian cases,” *IEEE Trans. Very Large Scale Integr. (VLSI) Syst.*, vol. 19, no. 9, pp. 1617–1626, Sep. 2011.
- [40] R. Cochran and S. Reda, “Spectral techniques for high-resolution thermal characterization with limited sensor data,” in *Proc. ACM 46th Annu. Design Autom. Conf.*, San Francisco, CA, USA, 2009, pp. 478–483.
- [41] M. V. Beigi and G. Memik, “THOR: Thermal-aware optimizations for extending ReRAM lifetime,” in *Proc. IEEE Int. Parallel Distrib. Process. Symp. (IPDPS)*, Vancouver, BC, Canada, 2018, pp. 670–679.
- [42] M. V. Beigi and G. Memik, “Thermal-aware optimizations of ReRAM-based neuromorphic computing systems,” in *Proc. ACM 55th Annu. Design Autom. Conf.*, San Francisco, CA, USA, 2018, p. 39.
- [43] W. Zhang and T. Li, “Helmet: A resistance drift resilient architecture for multi-level cell phase change memory system,” in *Proc. IEEE/IFIP 41st Int. Conf. Depend. Syst. Netw. (DSN)*, Hong Kong, 2011, pp. 197–208.
- [44] B. Datta and W. Burleson, “A high sensitivity and process tolerant digital thermal sensing scheme for 3-D ICS,” in *Proc. ACM 21st Ed. Great Lakes Symp. VLSI*, Lausanne, Switzerland, 2011, pp. 289–294.
- [45] D. Li, J.-H. Kim, and S. O. Memik, “Integrating thermocouple sensors into 3D ICS,” in *Proc. IEEE 31st Int. Conf. Comput. Design (ICCD)*, Asheville, NC, USA, 2013, pp. 221–226.
- [46] D. Li, S. Joshi, J.-H. Kim, and S. Ogrenci-Memik, “End-to-end analysis of integration for thermocouple-based sensors into 3-D ICS,” *IEEE Trans. Very Large Scale Integr. (VLSI) Syst.*, vol. 25, no. 9, pp. 2498–2511, Sep. 2017.
- [47] D. Jevdjic, S. Volos, and B. Falsafi, “Die-stacked DRAM caches for servers: Hit ratio, latency, or bandwidth? Have it all with footprint cache,” *ACM SIGARCH Comput. Archit. News*, vol. 41, no. 3, pp. 404–415, 2013.

- [48] T. E. Carlson, W. Heirman, and L. Eeckhout, "Sniper: Exploring the level of abstraction for scalable and accurate parallel multi-core simulations," in *Proc. Int. Conf. High Perform. Comput. Netw. Storage Anal. (SC)*, Nov. 2011, pp. 1–12.
- [49] Y. Kim, W. Yang, and O. Mutlu, "Ramulator: A fast and extensible DRAM simulator," *IEEE Comput. Archit. Lett.*, vol. 15, no. 1, pp. 45–49, Jan./Jun. 2016.
- [50] S. J. E. Wilton and N. P. Jouppi, "CACTI: An enhanced cache access and cycle time model," *IEEE J. Solid-State Circuits*, vol. 31, no. 5, pp. 677–688, May 1996.
- [51] S. Li, J. H. Ahn, R. D. Strong, J. B. Brockman, D. M. Tullsen, and N. P. Jouppi, "McPAT: An integrated power, area, and timing modeling framework for multicore and manycore architectures," in *Proc. 42nd Annu. IEEE/ACM Int. Symp. Microarchit. (MICRO)*, New York, NY, USA, 2009, pp. 469–480.
- [52] K. Skadron, M. R. Stan, K. Sankaranarayanan, W. Huang, S. Velusamy, and D. Tarjan, "Temperature-aware microarchitecture: Modeling and implementation," *ACM Trans. Archit. Code Optim.*, vol. 1, no. 1, pp. 94–125, 2004.
- [53] J. E. Miller *et al.*, "Graphite: A distributed parallel simulator for multicores," in *Proc. IEEE 16th Int. Symp. High Perform. Comput. Archit. (HPCA)*, Bengaluru, India, 2010, pp. 1–12.
- [54] *DDR4 SDRAM System-Power Calculator*. Accessed: Oct. 2018. [Online]. Available: https://www.micron.com/~media/documents/products/power-calculator/ddr4_power_calc.xlsm?la=en
- [55] W. Heirman, S. Sarkar, T. E. Carlson, I. Hur, and L. Eeckhout, "Power-aware multi-core simulation for early design stage hardware/software co-optimization," in *Proc. IEEE 21st Int. Conf. Parallel Archit. Compilation Tech. (PACT)*, Minneapolis, MN, USA, 2012, pp. 3–12.
- [56] J. Meng, K. Kawakami, and A. K. Coskun, "Optimizing energy efficiency of 3-D multicore systems with stacked DRAM under power and thermal constraints," in *Proc. ACM Annu. Design Autom. Conf.*, San Francisco, CA, USA, 2012, pp. 648–655.
- [57] *Spec cpu® 2006*. Accessed: Jan. 2019. [Online]. Available: <https://www.spec.org/cpu2006/>
- [58] *Intel Pentium M Processor 1.60 GHz, 1M Cache, 400 MHz FSB*. Accessed: Jan. 2019. [Online]. Available: https://ark.intel.com/products/27577/Intel-Pentium-M-Processor-1_60-GHz-1M-Cache-400-MHz-FSB
- [59] I. Bhati, M.-T. Chang, Z. Chishti, S.-L. Lu, and B. Jacob, "DRAM refresh mechanisms, penalties, and trade-offs," *IEEE Trans. Comput.*, vol. 65, no. 1, pp. 108–121, Jan. 2016.
- [60] A. Bilas, D. Jiang, and J. P. Singh, "Accelerating shared virtual memory via general-purpose network interface support," *ACM Trans. Comput. Syst.*, vol. 19, no. 1, pp. 1–35, Feb. 2001.
- [61] Y. Eckert, N. Jayasena, and G. H. Loh, "Thermal feasibility of die-stacked processing in memory," in *Proc. 2nd Workshop Near Data Process. (WoNDP)*, Cambridge, U.K., 2014, Art. no. 1.
- [62] Y. Zhu, B. Wang, D. Li, and J. Zhao, "Integrated thermal analysis for processing in die-stacking memory," in *Proc. ACM 2nd Int. Symp. Memory Syst.*, Alexandria, VA, USA, 2016, pp. 402–414.
- [63] K. Kang, G. De Micheli, S. Lee, and C.-M. Kyung, "Temperature-aware runtime power management for chip-multiprocessors with 3-D stacked cache," in *Proc. IEEE 15th Int. Symp. Qual. Electron. Design (ISQED)*, Santa Clara, CA, USA, pp. 163–170, 2014.
- [64] M. V. Beigi and G. Memik, "TAPAS: Temperature-aware adaptive placement for 3D stacked hybrid caches," in *Proc. ACM 2nd Int. Symp. Memory Syst.*, Alexandria, VA, USA, 2016, pp. 415–426.
- [65] D. Zhao, H. Homayoun, and A. V. Veidenbaum, "Temperature aware thread migration in 3D architecture with stacked DRAM," in *Proc. IEEE 14th Int. Symp. Qual. Electron. Design (ISQED)*, Santa Clara, CA, USA, 2013, pp. 80–87.
- [66] W.-H. Lo, K.-Z. Liang, and T. T. Hwang, "Thermal-aware dynamic page allocation policy by future access patterns for hybrid memory cube (HMC)," in *Proc. IEEE Design Autom. Test Europe Conf. Exhibit. (DATE)*, 2016, pp. 1084–1089.
- [67] D. Li, K. Zhang, A. Guliani, and S. Ogren-Memik, "Adaptive thermal management for 3D ICS with stacked DRAM caches," in *Proc. ACM 54th Annu. Design Autom. Conf.*, Austin, TX, USA, 2017, p. 3.
- [68] C. Yoon, J. H. Shim, B. Moon, and J. Kong, "3D die-stacked DRAM thermal management via task allocation and core pipeline control," *IEICE Electron. Exp.*, vol. 15, no. 3, 2018, Art. no. 20171253.
- [69] D. H. Albonesi, "Selective cache ways: On-demand cache resource allocation," in *Proc. IEEE 32nd Annu. Int. Symp. Microarchit. (MICRO)*, Haifa, Israel, 1999, pp. 248–259.



Minxuan Zhou received the B.S. degree in computer science and technology from Beihang University, Beijing, China, in 2015. He is currently pursuing the Ph.D. degree with the Department of Computer Science and Engineering, University of California at San Diego, San Diego, CA, USA.



Andreas Prodromou received the B.S. and M.Sc. degrees from the University of Cyprus, Nicosia, Cyprus, in 2011 and 2013, respectively. He is currently pursuing the Ph.D. degree with the University of California at San Diego, San Diego, CA, USA.



Rui Wang (M'10) received the B.S. and M.S. degrees in computer science from Xi'an Jiaotong University, Xi'an, China, in 2000 and 2003, respectively, and the Ph.D. degree in computer architecture from Beihang University, Beijing, China, in 2009.

He is an Assistant Professor with the School of Computer Science and Engineering, Beihang University. His current research interests include computer architecture and computer networks.

Dr. Wang is a member of China Computer Federation.



Hailong Yang received the Ph.D. degree from the School of Computer Science and Engineering, Beihang University, Beijing, China, in 2014.

He is an Assistant Professor with the School of Computer Science and Engineering, Beihang University. He has been involved in several scientific projects, such as performance analysis for big data systems and performance optimization for large scale applications. His current research interests include parallel and distributed computing, HPC, performance optimization, and energy efficiency.

Dr. Yang is a member of China Computer Federation.



Depei Qian received the master's degree from the University of North Texas, Denton, TX, USA, in 1984.

He is a Professor with the Department of Computer Science and Engineering, Beihang University, Beijing, China. He is currently serving as the Chief Scientist of China National High Technology Program on high productivity computer and service environment. His current research interests include innovative technologies in distributed computing, high-performance computing,

and computer architecture.

Prof. Qian is a fellow of China Computer Federation.



Dean Tullsen (F'07) received the B.S. and M.S. degrees from the University of California at Los Angeles, Los Angeles, CA, USA, in 1984 and 1986, respectively, and the Ph.D. degree from the University of Washington, Seattle, WA, USA, in 1996.

He is a Professor and the Chair of the Department of Computer Science and Engineering, University of California at San Diego, San Diego, CA, USA. He researches in the area of computer architecture, and also dabbles in compilers and security.

Prof. Tullsen was a recipient of the ISCA Influential Paper Award twice. He is a fellow of ACM and a past Chair of the IEEE Technical Committee on Computer Architecture.